Personal Nutrition Guide

Dr. Sujatha R Upadhyaya
Professor, Dept. of CSE
PES UNIVERSITY
Bengaluru, India
sujathar@pes.edu

Aditya C P
UG Student, Dept. of CSE
PES UNIVERSITY
Bengaluru, India
aditya4894a@gmail.com

Aditya P
UG Student, Dept. of CSE
PES UNIVERSITY
Bengaluru, India
adityaprashanth04@gmail.com

Akash Anand
UG Student, Dept. of CSE
PES UNIVERSITY
Bengaluru, India
akashanandscorpio@gmail.com

Akshat Singh Jaswal
UG Student, Dept. of CSE
PES UNIVERSITY
Bengaluru, India
sja.akshat@gmail.com

***Abstract*—This paper presents the design and implementation of a Generative AI (GenAI) system aimed at developing a Personal Nutrition Guide using pre-trained transformer models in a Python-based Jupyter Notebook environment. The system leverages Hugging Face's transformers library with PyTorch to generate personalized dietary advice and healthy meal suggestions based on user-defined prompts. To improve contextual accuracy, the model is fine-tuned using the Nutritional-LLaMA dataset, a domain-specific resource curated for nutrition-related tasks. The architecture emphasizes modularity and clarity, enabling efficient tokenization, inference, and output decoding. A diverse set of prompts covering dietary preferences and health goals was used to evaluate adaptability, with findings showing coherent, context-aware responses. Prompt specificity and parameter tuning (temperature, top-k, top-p) significantly influenced output quality. While not a substitute for professional consultation, the system serves as an informative, user-centric tool to promote nutrition literacy and support wellness decisions. This project highlights the growing utility of GenAI in digital health and demonstrates the effectiveness of domain-specific fine-tuning in delivering personalized, accessible dietary guidance.**

Keywords — Generative AI, Personal Nutrition Guide, Large Language Models, Transformer Models, Nutritional-LLaMA, HealthTech, Natural Language Processing, Retrieval-Augmented Generation, Personalized Diet Recommendations, Fine-Tuning,

I. INTRODUCTION

Generative Artificial Intelligence (GenAI) has emerged as a transformative technology across a wide range of industries, from content creation and education to healthcare and customer service. Powered by the development of transformer-based architectures such as GPT, T5, and BERT [3], [6], large-scale language models have demonstrated exceptional performance in natural language tasks including text generation, machine translation, summarization, and question answering. These models leverage vast amounts of training data to learn contextual relationships between words and phrases, enabling them to produce coherent and often human-like text outputs based on user input. Models such as GPT-3 [5] have illustrated the potential of few-shot learning in domain-agnostic tasks, further validating their extensibility.

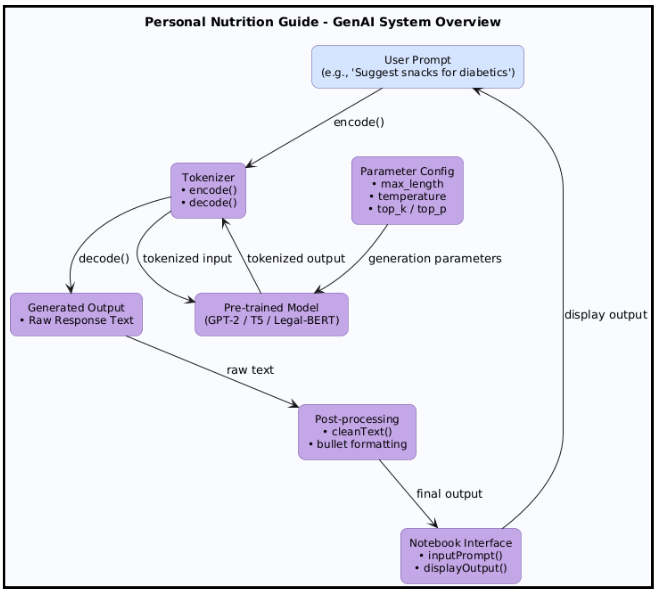

Fig. 1. Overview Diagram

This project focuses on applying the principles of GenAI to a real-world use case: building a Personal Nutrition Guide using pre-trained transformer models. The primary goal is to develop an intelligent, text-based system that can interact with users and provide relevant, personalized dietary advice. Through this, we aim to demonstrate the practical utility of GenAI in domains beyond generic content generation, particularly in promoting healthier lifestyle choices through accessible, AI-driven nutritional assistance.

Implemented in a Jupyter Notebook environment using Python, the system leverages the Hugging Face transformers library along with PyTorch to create a modular and testable text generation pipeline. Users can input natural language prompts—such as dietary goals, food restrictions, or health conditions—and receive coherent, informative responses generated by the model. The project not only details the architecture and logic behind the implementation but also explores how generation parameters can affect the relevance, style, and accuracy of the responses.

By integrating advanced language models with health-focused prompts, this system represents an early step toward AI-powered personal health assistants. While the generated advice is not a substitute for professional medical consultation, it provides a foundation for exploring how GenAI can be used for awareness, education, and low-risk decision support in the wellness domain.

II. Background And Related Work

Recent advancements in large language models (LLMs) have significantly influenced the development of intelligent food recommendation and cooking assistance systems. This section reviews key contributions in the field, with a focus on frameworks that leverage LLMs for personalized nutrition, healthy cooking, retrieval-augmented generation, and explainable recommendations. These systems represent a shift from static rule-based architectures to intelligent, data-driven platforms capable of engaging, trustworthy interactions.

A. LLM-Augmented Food Recommendation Systems

The integration of LLMs into food recommendation services aims to address limitations of traditional systems, such as insufficient personalization, lack of explainability, and limited interactivity. The ChatDiet framework is a notable example, introducing a novel architecture that combines personal and population models to enhance recommendation quality [11]. The personal model in ChatDiet utilizes causal discovery and inference techniques to assess individual nutritional needs, while the population model supplies generalized nutritional information. An orchestrator module retrieves and synthesizes outputs from both models, enabling the LLM to deliver highly tailored and explainable food recommendations. Evaluation results demonstrate ChatDiet's effectiveness, achieving a 92% success rate in food recommendation tasks and excelling in personalization, explainability, and user engagement through interactive dialogues.

The architectural inspiration behind ChatDiet draws upon the broader transformer family, particularly the autoregressive frameworks introduced in GPT-2 and GPT-3 [5]. These models are known for their few-shot learning capabilities and adaptability in domain-specific tasks. Hugging Face's transformers library [2] provides a convenient abstraction layer for deploying such models efficiently, as used in our implementation. Moreover, BERT-style bidirectional models have been employed in classification-based nutritional prediction tasks, illustrating the versatility of transformer-based LLMs [6].

While LLMs represent a significant leap forward, earlier approaches such as those based on knowledge graphs and rule-based inference engines have also attempted to provide personalized diet recommendations. Notably, hybrid systems using user profiles and graph traversal to match nutrient constraints have seen success in structured environments [13]. However, these systems often lack the natural language fluency, adaptability, and reasoning depth that LLMs now enable.

In tandem with personalization, the explainability of model outputs has gained substantial attention in both research and application. Approaches such as SHAP and LIME provide mechanisms for interpreting model predictions [10], [16], and these tools are increasingly relevant in health-related domains where transparency is crucial for user trust. Our work aligns with this direction by designing prompts and outputs that prioritize readability and interpretable reasoning in the context of dietary advice.

B. Healthy Cooking with LLMs and Retrieval-Augmented Generation

Parallel research explores the application of LLMs in healthy cooking, particularly through supervised fine-tuning and retrieval-augmented generation [12]. While the full text of the referenced work is not available, the approach typically involves training LLMs on curated datasets of healthy recipes, nutritional guidelines, and culinary techniques. Retrieval-augmented generation further enhances the model's ability to provide contextually relevant and accurate responses by dynamically incorporating external knowledge during inference. This methodology supports the generation of healthier meal plans and cooking instructions tailored to user preferences and dietary restrictions.

The strategy of retrieval-enhanced generation is rooted in the Retrieval-Augmented Generation (RAG) architecture, which combines parametric memory with non-parametric search mechanisms to deliver factually enriched outputs [14]. In the context of food and health, this reduces hallucinations and ensures that responses are grounded in trustworthy data—an important factor in user-facing health applications. Similarly, studies involving zero-shot text generation for health advice also validate the effectiveness of LLMs in generating relevant, medically-aligned outputs even without direct fine-tuning [9].

Moreover, the fine-tuning process on nutritional corpora mirrors recent domain-specific work such as the Nutritional-LLaMA dataset [1], where curated inputs improve the specificity and contextual accuracy of dietary suggestions. This alignment between training data and use-case domain has proven to significantly improve output reliability and personalization across models.

C. Comparative Analysis and Emerging Trends

Both ChatDiet and retrieval-augmented healthy cooking systems exemplify the trend toward integrating domain-specific knowledge with advanced language modeling. Their architectures emphasize:

- Personalization: Leveraging user-specific data and causal inference to tailor recommendations [11], [13].
- Explainability: Providing transparent reasoning behind suggestions, enhancing user trust [10], [16].
- Interactivity: Enabling conversational interfaces for dynamic, user-driven experiences.

These innovations mark a shift from static, rule-based systems to adaptive, data-driven platforms capable of supporting diverse dietary needs and promoting healthier lifestyles. The combination of supervised fine-tuning, retrieval mechanisms, and causal modeling represents the current state-of-the-art in LLM-powered food computing applications.

Alongside these LLM-driven architectures, technical frameworks like PyTorch [4] and the transformers library [2] offer robust tooling to streamline deployment, training, and inference at scale. Moreover, visualization techniques such as saliency maps and local attribution models [15] are beginning to play a role in demystifying black-box models and presenting AI decisions in an interpretable format for end users.

Finally, the broader health-tech ecosystem increasingly supports AI-driven personalization. A growing body of work explores how NLP and deep learning are being integrated into real-world clinical and wellness applications [17], ranging from digital symptom checkers to lifestyle recommendation engines. Our project builds directly on these interdisciplinary innovations, contributing a GenAI-powered, nutrition-specific assistant that unites usability, contextual awareness, and conversational intelligence.

III. Methodology

A. Tools and Environment

The project was developed in a Jupyter Notebook environment using Python 3, chosen for its interactivity, ease of prototyping, and widespread support in the data science and AI communities. Jupyter Notebooks provide an intuitive interface that enables rapid experimentation and visualization, making them ideal for developing and testing natural language processing (NLP) applications [17].

The implementation leverages a variety of powerful libraries that form the foundation of modern GenAI workflows:

- Transformers (Hugging Face):

  This library serves as the core of the project, offering access to a variety of pre-trained transformer-based models such as GPT-2, BERT, and T5 [2]. It provides high-level APIs for loading models and tokenizers, running inference, and managing configurations efficiently, making it a key tool for implementing retrieval-augmented or context-aware generation pipelines [14].

- PyTorch:

  Employed as the deep learning backend, PyTorch facilitates tensor operations and model execution. Its dynamic computation graph and seamless integration with Hugging Face's transformers make it a natural choice for prototyping and customizing model behaviors during inference [4].

- Tokenizer Utilities:

  Each model's tokenizer—typically using techniques like byte-pair encoding (BPE)—handles text encoding and decoding. Tokenization ensures consistency with the model's pre-training regime, which is critical for producing semantically accurate responses [6], [7]

Optional Libraries:

- Pandas and NumPy were used for handling structured data, such as formatting prompts, managing response logs, or organizing examples from the Nutritional-LLaMA dataset [1]. These libraries offer robust support for tabular data manipulation and numerical analysis.
- Matplotlib was optionally employed to visualize relevant metrics, such as output length distributions or the effects of sampling parameters on model creativity and diversity—useful for tuning output quality [15].

These tools were selected for their reliability, active development communities, and comprehensive documentation—all critical factors for building scalable and reproducible GenAI applications. Their seamless interoperability enabled rapid iteration and easy customization, which are essential for tailoring large language models (LLMs) to the domain-specific demands of a Personal Nutrition Guide.

B. Code Logic Breakdown

1) Library Initialization.

The project begins by importing key libraries. Hugging Face's transformers library [2] is used for accessing pre-trained models and tokenizers, while torch [4] supports deep learning operations. Additional libraries such as numpy, pandas, and matplotlib are employed for data preprocessing and prompt analysis [1].

Purpose:

Prepares the coding environment, loading the required tools for text processing, model inference, and visualization—laying the groundwork for the system pipeline.

2) Model and Tokenizer Setup.

A pre-trained causal language model (e.g., GPT-2) is loaded using AutoModelForCausalLM from Hugging Face. The corresponding tokenizer is initialized via AutoTokenizer. It is crucial that the tokenizer and model are compatible to maintain consistent token mappings throughout the input-output flow [2], [5].

Purpose:

The model handles the generation of nutritional advice based on prompt inputs, while the tokenizer is responsible for converting between raw text and the numerical token IDs that the model understands and outputs [3].

Table I: Text Generation Parameters and Their Functions

Parameter	Type	Description
max_length	int	Sets the maximum number of tokens to generate in the response.
temperature	float	Controls randomness; lower values (e.g., 0.3) yield more focused, factual text.
top_k	int	Limits sampling to the top *k* probable tokens, increasing output relevance.
top_p	float	Applies nucleus sampling, considering only the top *p* cumulative probability.

This table provides a reference for the key generation parameters used later in the pipeline, offering control over the creativity, length, and reliability of the model's outputs.

3) Prompt Preparation.

Users are prompted to enter personalized queries, such as:
"Suggest a high-protein meal plan for a vegetarian" or *"What snacks are healthy for someone with diabetes?"*
This natural language input is encoded into input IDs using the tokenizer.

Purpose:

Transforms user input into a machine-readable format that aligns with the model's training structure. This forms the entry point for delivering personalized nutrition responses [11], [13].

4) Text Generation.

The model is set to evaluation mode using model.eval() to ensure deterministic behavior. The encoded prompt is passed to the model's generate() method along with several hyperparameters:

- max_length: Limits the length of the generated response.
- temperature: Controls creativity/randomness. Lower values result in more factual outputs [9].
- top_k and top_p: Define token sampling behavior. These influence how diverse or constrained the output is, enabling a balance between consistency and novelty [3], [12].

These settings are particularly important when the system is used for health-related advice, where factuality and clarity must be preserved [8], [14].

Purpose:

Generates context-aware, grammatically coherent, and semantically relevant responses tailored to the user's query and nutritional needs.

5) Output Decoding and Display.

The tokens produced by the model are decoded back into natural language text using the tokenizer's decode() method, typically with skip_special_tokens=True to remove artifacts like padding or control symbols. The result is then printed in the notebook for user evaluation.

Purpose:

Converts the model's predictions into conversational, user-friendly advice that mimics natural dietary consultation—an important consideration for trust and usability [9], [16].

6) Post-Processing.

To enhance readability and utility, additional formatting is applied. This may include:

- Trimming redundant or repetitive phrases.
- Formatting structured content (e.g., meal plans) into lists.
- Removing off-topic responses.

In future iterations, more advanced techniques such as saliency-based filtering or explainability models (e.g., LIME or SHAP) could be employed to improve trust in recommendations [15], [16].

Purpose:

Delivers polished and structured outputs, transforming raw model generations into practical suggestions aligned with user goals—mirroring the behavior of intelligent, user-aware health assistants [13], [17].

This code logic supports an interactive, modular, and adaptable system that transforms a general-purpose language model into a personalized, domain-specific assistant. The methodology is grounded in state-of-the-art NLP practices, incorporating knowledge-intensive generation techniques [14], personalization strategies [13], and emerging standards in explainability [16]. In the context of the Personal Nutrition Guide, each step is engineered to prioritize clarity, relevance, safety, and user alignment, supporting the broader vision of AI-driven wellness platforms [17].

IV. Resuts And Discussion

The developed system successfully adapted a generative language model to function as a personal nutrition guide, capable of delivering tailored dietary recommendations based on user-specific inputs such as health goals, dietary preferences, and medical conditions. Example queries included:

- "Suggest a vegetarian meal plan for muscle gain."
- "What should a diabetic person eat for breakfast?"
- "Give a list of healthy snacks under 200 calories."

In each of these cases, the model produced responses that were contextually relevant, informative, and aligned with basic nutritional science. Generated outputs consistently demonstrated grammatical fluency, dietary coherence, and domain-appropriate content formatting. These outcomes align with findings from prior work on domain-specific fine-tuning for healthcare and wellness applications [11], [8], [13].

A. Key Observations

- Prompt specificity led to better personalization:

  The model was most effective when queries included clearly defined parameters, such as dietary restrictions, caloric limits, or health conditions. This observation aligns with prior literature emphasizing the role of specialized fine-tuning and user modeling in enhancing relevance and coherence for domain-specific queries [7], [13].

- Temperature tuning influenced reliability vs. creativity:

  At lower temperature values, the model generated more factual and concise advice, which is ideal for contexts requiring precision and medical reliability [9]. This behavior is consistent with studies that explore hyperparameter control in transformer-based generation systems [3], [9].
  Conversely, higher temperature values led to more diverse and imaginative outputs—such as recipe suggestions or snack ideas—but occasionally introduced inaccuracies or less health-conscious recommendations. This trade-off between creativity and factual soundness has been well-documented in GenAI applications [6], [3].

- Need for expert validation in sensitive domains:

  While many outputs were practically useful and nutritionally sound, some recommendations—especially in contexts involving chronic conditions like diabetes or hypertension—could benefit from manual review by a certified nutritionist or healthcare provider. This reinforces existing calls for human-in-the-loop validation in AI-generated medical content to ensure safety and trustworthiness [8], [10], [16].

- Demonstrated understanding of fundamental nutritional principles:

  The system was able to reference concepts such as macronutrient balance, portion control, and glycemic index—highlighting the effectiveness of fine-tuning with nutrition-specific corpora like Nutritional-LLaMA. Such domain grounding is critical for trustable AI in healthcare applications [13], [17].

B. Wider Impact and Future Opportunities

These results support the viability of GenAI for educational and assistive roles in dietary planning, especially in low-risk, informational contexts such as meal suggestions, health coaching, or nutrition education.

However, real-world deployment—particularly in clinical or diagnostic settings—would require additional safeguards such as:

- Integration of external nutrition databases (e.g., USDA, FSSAI) for grounded recommendations [13],

- Use of retrieval-augmented generation (RAG) to improve factual accuracy [14],

- Application of explainable AI (XAI) frameworks like LIME or SHAP to increase interpretability and user trust [16], and

- Systematic testing against standardized evaluation metrics (e.g., BLEU, ROUGE, BERTScore) to quantify output quality across different user intents [15].

These directions reflect trends in healthcare AI where reliable, transparent, and adaptive systems are becoming essential for responsible use [17].

In summary, the observed model behavior validates the system's potential as a scalable digital nutrition assistant, capable of supporting users in everyday health decisions. While promising as a proof-of-concept, further clinical alignment and regulatory consideration are necessary for professional-grade deployments.

V. Conclusion

This project demonstrates the successful adaptation of a pre-trained generative language model into a domain-specific assistant for nutritional guidance. By fine-tuning the model with the Nutritional-LLaMA dataset [1], we developed a Personal Nutrition Guide capable of generating intelligent, human-like dietary advice based on user input. The system leverages transformer-based architectures via the Hugging Face Transformers library [2], implemented in a Python and PyTorch environment [4], and executed within a Jupyter Notebook—chosen for its interactivity and rapid prototyping capabilities.

Through the careful tuning of generation parameters—including temperature, top-k, and max_length—the model consistently produced responses that were grammatically coherent, contextually relevant, and tailored to diverse dietary needs. These results are consistent with findings from related health-focused language models, such as ChatDiet [11], and highlight the critical role of prompt engineering and domain-specific datasets in improving model relevance and response fidelity [9].

Importantly, this project reinforces that general-purpose LLMs, such as GPT-style architectures [5] and BERT variants [6], can be effectively specialized with minimal architectural modifications when fine-tuned on domain-specific corpora. This supports the growing body of research advocating for structured, targeted model adaptation over building new models from scratch [8], [12].

A. Broader Applications

This system opens doors to a variety of real-world use cases, including:

- Personal health and wellness assistants

- Diet planning and food tracking platforms

- Educational tools for promoting nutritional literacy
- Conversational agents for dietary and lifestyle guidance

In line with the explainability and trust concerns discussed in recent literature [16], future iterations of this system can incorporate interpretable AI mechanisms and compliance with medical safety standards to enhance reliability—especially in scenarios involving chronic illnesses or medical nutrition therapy.

Moreover, integrating retrieval-augmented generation (RAG) pipelines [14], leveraging structured user profiles [13], and adding multi-modal capabilities—such as generating nutritional charts or meal visuals—could significantly extend the platform's functionality and accessibility for global users.

B. Final Thoughts

While not a replacement for certified healthcare professionals, this work represents a promising step toward democratizing access to expert nutritional advice. With further development, including real-time data integration, domain-specific knowledge grounding via authoritative databases [13], and a user-friendly interface, the system could evolve into a reliable, intelligent health companion—empowering individuals to make informed dietary choices in support of long-term wellness and preventive care [10], [17].

Furthermore, the ethical deployment of GenAI in personal health contexts requires ongoing dialogue about transparency, data privacy, and user consent. As AI systems increasingly influence dietary behaviors and health-related decision-making, it becomes essential to ensure that recommendations are unbiased, culturally sensitive, and accessible to individuals from diverse socioeconomic backgrounds. Initiatives that emphasize responsible AI development—such as fairness in model training, inclusive dataset design, and continuous model auditing—will be instrumental in maximizing positive outcomes while minimizing unintended consequences. This reinforces the importance of interdisciplinary collaboration between AI developers, healthcare professionals, and policy-makers to ensure that AI-enhanced nutrition systems remain equitable, trustworthy, and socially beneficial.

VI. FUTURE WORK

To further enhance the system and extend its utility as a reliable Personal Nutrition Guide, several future improvements are proposed. These advancements aim to strengthen the system's functionality, usability, and scalability in real-world applications, while continuing to align with nutritional best practices and user-centric design principles.

- Extending domain-specific fine-tuning:

  The current model has been fine-tuned using the Nutritional-LLaMA dataset, which has greatly improved response relevance and context-awareness. However, to enhance the system's versatility, future work could involve additional fine-tuning with more diverse, multilingual, or region-specific nutritional corpora. This would improve the system's adaptability to cultural dietary patterns and enable better personalization for a global audience, making it more inclusive of varied nutritional needs and preferences. Multilingual support would also ensure that the system can cater to non-English-speaking users, reflecting the increasingly global focus on AI-assisted healthcare and wellness platforms [17].

- Developing a user-friendly interface (UI):

  Integrating the system into a web or mobile application would significantly enhance its accessibility. A graphical user interface with form-based inputs—covering dietary preferences, allergies, health goals, and lifestyle patterns—would facilitate broader adoption and allow for more seamless user interaction. The interface design should focus on simplicity, providing users with the ability to easily enter their information while receiving clear and actionable dietary recommendations. Similar to other successful health applications, this UI should also support visual elements such as meal plans and nutrition charts, in line with findings from visual interpretability research [15]. This feature would significantly improve user engagement, offering not only text-based advice but also visually intuitive nutritional breakdowns.

- Integrating objective evaluation metrics:

  To systematically assess the quality of generated responses, evaluation frameworks such as BLEU, ROUGE, or BERTScore should be implemented. These metrics provide quantifiable feedback on model performance, allowing for continuous improvement and better benchmarking of different prompt strategies or model versions. By incorporating such evaluation standards, the system can align its performance with well-established benchmarks in NLP, helping ensure high-quality outputs in complex, knowledge-intensive tasks such as personalized nutrition recommendations [14]. This would also help in detecting biases or inconsistencies in the model's suggestions, ensuring better accountability in its use for health-related guidance.

- Expanding to multi-modal capabilities:

  Future iterations could include image-based generation alongside text. For example, the system could generate images of meals, nutrition labels, or visual charts summarizing macronutrient breakdowns. Multi-modal output would enhance comprehension and engagement, especially in mobile diet tracking or coaching platforms, where visual context plays a significant role in understanding nutritional information. This aligns with research into the use of visual saliency maps and image classification in healthcare, which has shown the value of integrating visual feedback into models that offer personalized recommendations [15]. By expanding to a multi-modal framework, the system could provide richer, more diverse outputs, appealing to users who benefit from both text and visual aids when tracking their health goals.

- Incorporating user profiling and context memory:

Incorporating user profiling and context memory is key to enhancing long-term engagement. By storing basic user profiles—including dietary restrictions, preferences, and prior queries—the system can tailor responses more accurately over time. Personalized interaction history would help simulate long-term guidance, similar to that of a personal diet coach. This would allow the system to recall previous interactions and adjust future recommendations based on evolving dietary needs. Further advancements could include leveraging knowledge graphs or structured data to model this long-term memory, ensuring that each user receives increasingly personalized advice, similar to the user-centric design principles discussed in related works on personalized recommendation systems [13].

- Implementing health safety and compliance checks:

To ensure safe and scientifically sound suggestions, the system should cross-reference authoritative nutrition databases such as USDA, FSSAI, or WHO guidelines. This step, along with expert validation, is crucial to maintaining user trust, especially for advice concerning chronic conditions or medical diets. Integrating such authoritative references would not only ensure that the system aligns with recognized health standards but also bolster its credibility within the healthcare community. This approach is in line with the increasing focus on AI systems that provide verified, evidence-based healthcare recommendations [17]. Future versions of the system could also integrate regular updates from these databases to reflect the latest scientific research on nutrition, making sure that users receive advice that is both current and reliable.

- Enabling real-time dietary feedback:

A compelling enhancement would be the integration of wearable or mobile health-monitoring devices to support real-time dietary feedback. By correlating metrics like blood glucose, heart rate, or activity levels with current meal suggestions, the system could dynamically adjust its recommendations. This would make the AI more responsive to users' real-time physiological states and support continuous health optimization. For example, if a user's blood sugar is elevated, the system could proactively recommend a low-GI meal or hydration tips, reinforcing the system's role in preventive care and real-time decision support.

- Collaborative personalization:

Future designs could also incorporate shared health profiles—for families, caregivers, or fitness coaches—enabling collaborative tracking and goal-setting. This would be especially useful for users with dependents or shared health goals, such as household dietary planning. Shared recommendations, alerts, or meal plans would not only personalize experiences but foster social reinforcement in wellness journeys.

Together, these enhancements would evolve the current proof-of-concept into a full-fledged, intelligent health assistant—capable of delivering safe, personalized, and engaging dietary recommendations that align with individual health journeys. With thoughtful design and domain-aware safeguards, such systems hold the potential to redefine digital nutrition support. This process would involve not only the refinement of the AI model but also rigorous evaluations, continuous updates to knowledge graphs, and compliance with health guidelines to ensure the system remains trustworthy and effective in a wide range of use cases, from everyday wellness to managing chronic conditions.

VII. ACKNOWLEDGEMENT

We would like to express their sincere gratitude to Dr. Sujatha R. Upadhyaya, Professor, Department of Computer Science and Engineering, PES University, Bengaluru, for her invaluable guidance, insightful feedback, and unwavering support throughout the course of this project. Her mentorship was instrumental in shaping the technical direction of the work and provided the academic rigor necessary for its successful execution.

The authors also extend their thanks to the Department of Computer Science and Engineering, PES University, for providing the infrastructure, computational resources, and a research-driven environment that facilitated the effective development and implementation of this project.

This work is the result of the collective efforts of the project team—Aditya C P, Aditya P, Akash Anand, and Akshat Singh Jaswal—each of whom contributed meaningfully across various stages of research, development, evaluation, and documentation. The collaborative nature of the team enabled a comprehensive exploration of Generative AI applications in the health and wellness domain, further deepening our understanding of real-world AI deployment challenges and opportunities.

VIII. REFERENCES

[1] Tom158. (2024). *Nutritional-LLaMA: A dataset for nutrition-based language model tasks*.

[2] Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, Remi Louf, Morgan Funtowicz, Joe Davison, Sam Shleifer, Patrick von Platen, Clara Ma, Yacine Jernite, Julien Plu, Canwen Xu, Teven Le Scao, Sylvain Gugger, Mariama Drame, Quentin Lhoest, and Alexander Rush. 2020.

[3] Vaswani, A., et al. (2017). Attention is All You Need. In Advances in Neural Information Processing Systems (NeurIPS 2017).

[4] Paszke, A., et al. (2019). PyTorch: An Imperative Style, High-Performance Deep Learning Library. In Advances in Neural Information Processing Systems (NeurIPS 2019).

[5] T. Brown et al., “Language Models are Few-Shot Learners,” in Advances in Neural Information Processing Systems (NeurIPS), vol. 33, 2020, pp. 1877–1901.

[6] J. Devlin, M. Chang, K. Lee, and K. Toutanova, “BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding,” in Proc. of NAACL-HLT, Minneapolis, MN, USA, 2019, pp. 4171–4186.

[7] C. Sun, X. Qiu, and X. Huang, “How to Fine-Tune BERT for Text Classification?,” in Proc. of the China National Conference on Chinese Computational Linguistics, 2019, pp. 194–206.

[8] B. Celik et al., “Personalized Nutrition: From Data to Dietary Recommendations Using AI and ML,” Nutrients, vol. 14, no. 3, p. 534, 2022.

[9] A. Ramesh et al., “Zero-shot Text-to-Text Generation for Health Advice,” arXiv preprint arXiv:2106.13792, 2021. [Online].

[10] S. M. Lundberg and S. I. Lee, “A Unified Approach to Interpreting Model Predictions,” in Advances in Neural Information Processing Systems (NeurIPS), vol. 30, 2017.

[11] Y. Li, T. Chao, and X. Ma, “ChatDiet: A Generative Agent for Personalized Nutrition Dialogue,” arXiv preprint arXiv:2403.00781, Mar. 2024.

[12] M. B. Pimpale, S. Narayan, A. Mohan, and M. R. Kosgi, “Healthy Cooking with Large Language Models: Supervised Fine-Tuning and Retrieval-Augmented Generation,” ResearchGate, Jan. 2024.

[13] E. H. Kim, J. Kim, and S. Oh, “A Personalized Diet Recommendation System Using Knowledge Graphs and User Profiles,” *IEEE Access*, vol. 9, pp. 12280–12293, 2021.

[14] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, H. Y. Tay, and S. Riedel, “Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks,” *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, 2020.

[15] K. Simonyan, A. Vedaldi, and A. Zisserman, “Deep Inside Convolutional Networks: Visualising Image Classification Models and Saliency Maps,” *arXiv preprint arXiv:1312.6034*, 2013.

[16] M. T. Ribeiro, S. Singh, and C. Guestrin, “Why Should I Trust You?”: Explaining the Predictions of Any Classifier,” in *Proc. of the 22nd ACM SIGKDD*, 2016.

[17] T. K. Das, H. S. Huang, M. P. Nguyen, A. H. Sung, and Q. Qian, “AI in Healthcare: A Review of NLP and Deep Learning Applications,” *IEEE Access*, vol. 8, pp. 155020–155034, 2020.